

Theoretical Foundations of Distributed Operating Systems

Neeraj Mittal

The University of Texas at Dallas

June 1, 2012

Overview

Ordering Events

Abstract Clocks

Ordering of Messages

State of a Distributed System

Monitoring a Distributed System

Outline

Overview

Ordering Events

Abstract Clocks

Ordering of Messages

State of a Distributed System

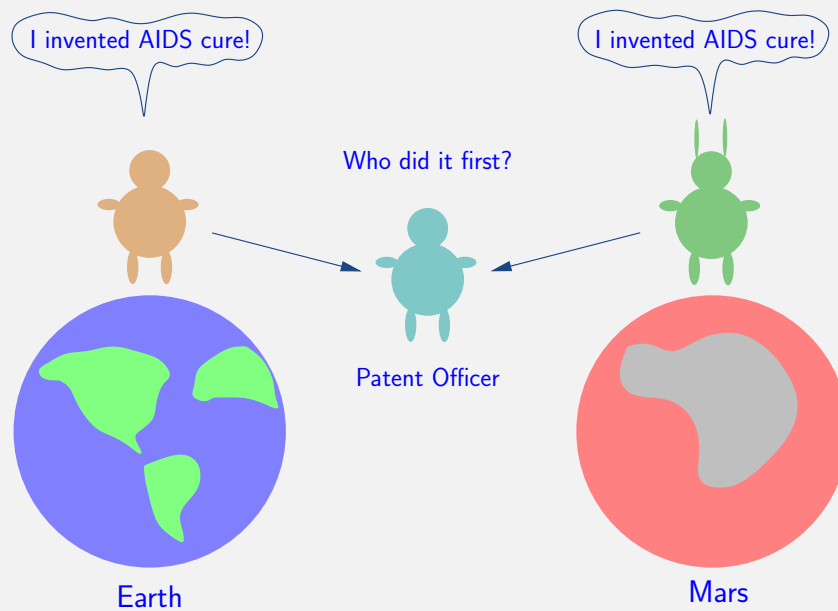
Monitoring a Distributed System

Two Important Characteristics

- ▶ Absence of Global Clock
 - ▶ there is **no common** notion of **time**
- ▶ Absence of Shared Memory
 - ▶ **no** process has **up-to-date knowledge** about the system

Absence of Global Clock


- Different processes may have **different notions of time**



Problem: How do we **order events** on different processes?

Absence of Shared Memory

- ▶ A process does not know **current state of other processes**



Problem: How do we obtain a **coherent view** of the system?

When is it possible to order two events?

► Three cases:

1. Events executed on the **same process**:

- if e and f are events on the same process and e occurred before f , then e *happened-before* f

2. Communication events of the **same message**:

- if e is the send event of a message and f is the receive event of the same message, then e *happened-before* f

3. Events related by **transitivity**:

- if event e happened-before event g and event g happened-before event f , then e *happened-before* f

Outline

Overview

Ordering Events

Abstract Clocks


Ordering of Messages

State of a Distributed System

Monitoring a Distributed System

Happened-Before Relation


- ▶ Happened-before relation is denoted by \rightarrow
- ▶ Illustration:



- ▶ Events on the same process:
examples: $a \rightarrow b$, $a \rightarrow c$, $d \rightarrow f$
- ▶ Events of the same message:
examples: $b \rightarrow e$, $f \rightarrow i$
- ▶ Transitivity:
examples: $a \rightarrow e$, $a \rightarrow i$, $e \rightarrow i$

Concurrent Events

- ▶ Events **not related** by happened-before relation
- ▶ Concurrency relation is denoted by \parallel
- ▶ Illustration:



- ▶ Examples: $a \parallel d$, $d \parallel h$, $c \parallel e$
- ▶ Concurrency relation is **not transitive**:
example: $a \parallel d$ and $d \parallel c$ but $a \nparallel c$

Outline

Overview

Ordering Events

Abstract Clocks

Ordering of Messages

State of a Distributed System

Monitoring a Distributed System

Different Kinds of Clocks

- ▶ Logical Clocks
 - ▶ used to **totally order** all events
- ▶ Vector Clocks
 - ▶ used to track **happened-before** relation
- ▶ Matrix Clocks
 - ▶ used to track what **other processes know** about other processes
- ▶ Direct Dependency Clocks
 - ▶ used to track **direct** causal dependencies

Logical Clock

- ▶ Implements the notion of **virtual time**
- ▶ Can be used to **totally order** all events
- ▶ Assigns timestamp to each event in a way that is **consistent with the happened-before** relation:

$$e \rightarrow f \implies C(e) < C(f)$$


$C(e)$: timestamp for event e

$C(f)$: timestamp for event f

Implementing Logical Clock

- ▶ Each process has a local **scalar** clock, initialized to zero
 - ▶ C_i denotes the local clock of process P_i
- ▶ Action depends on the type of the event
- ▶ Protocol for process P_i :
 - ▶ On executing an **interval event**:
 $C_i := C_i + 1$
 - ▶ On **sending a message** m :
 $C_i := C_i + 1$
piggyback C_i on m
 - ▶ On **receiving a message** m :
let t_m be the timestamp piggybacked on m
 $C_i := \max\{C_i, t_m\} + 1$

Implementing Logical Clock: An Illustration

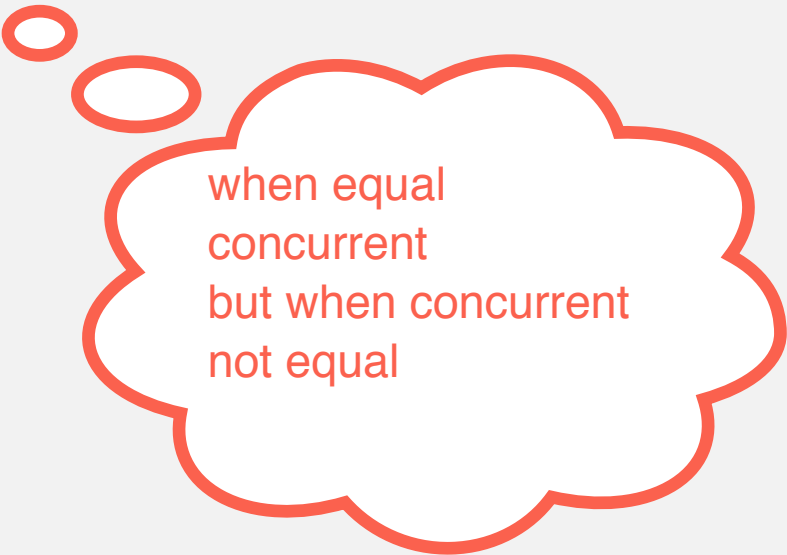


Limitation of Logical Clock

- ▶ Logical clock

cannot be used to determine whether two events are concurrent

$e \parallel f$ does not imply $C(e) = C(f)$



when equal
concurrent
but when concurrent
not equal

Vector Clock



- ▶ Captures the **happened-before** relation
- ▶ Assigns timestamp to each event such that:

$$e \rightarrow f \iff C(e) < C(f)$$

$C(e)$: timestamp for event e

$C(f)$: timestamp for event f

Comparing Two Vectors

- ▶ Vectors are compared **component-wise**:

- ▶ Equality:

$$V = V' \quad \text{iff} \quad \langle \forall i : V[i] = V'[i] \rangle$$

- ▶ Less Than:

$$V < V' \quad \text{iff} \quad \langle \forall i : V[i] \leq V'[i] \rangle \wedge \langle \exists i : V[i] < V'[i] \rangle$$

- ▶ Example:

$$\begin{bmatrix} 1 \\ 2 \\ 0 \end{bmatrix} < \begin{bmatrix} 2 \\ 3 \\ 1 \end{bmatrix} \quad \text{and} \quad \begin{bmatrix} 2 \\ 1 \\ 1 \end{bmatrix} < \begin{bmatrix} 2 \\ 3 \\ 4 \end{bmatrix} \quad \text{but} \quad \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \not< \begin{bmatrix} 2 \\ 1 \\ 3 \end{bmatrix}$$

Implementing Vector Clock


- ▶ Each process has a local **vector** clock
 - ▶ C_i denotes the local clock of process P_i
- ▶ Action depends on the type of the event
- ▶ Protocol for process P_i :
 - ▶ On executing an **interval event**:
$$C_i[i] := C_i[i] + 1$$
 - ▶ On **sending a message** m :
$$C_i[i] := C_i[i] + 1$$

piggyback C_i on m
 - ▶ On **receiving a message** m :

let t_m be the timestamp piggybacked on m

$$\forall k \quad C_i[k] := \max\{C_i[k], t_m[k]\}$$
$$C_i[i] := C_i[i] + 1$$

Implementing Vector Clock: An Illustration



Properties of Vector Clock

- ▶ How many **comparisons** are needed to determine whether an event e happened-before another event f ?
 - ▶ As many as N integers may need to be compared in the worst case, where N is the number of processes
 - ▶ Can we do better?
 - ▶ Suppose we know the processes on which events e and f occurred (say, P_i and P_j respectively)

$$e \rightarrow f$$

if and only if

$$(i = j) \wedge (C(e)[i] < C(f)[i]) \vee (i \neq j) \wedge (C(e)[i] \leq C(f)[i])$$

Outline

Overview

Ordering Events

Abstract Clocks

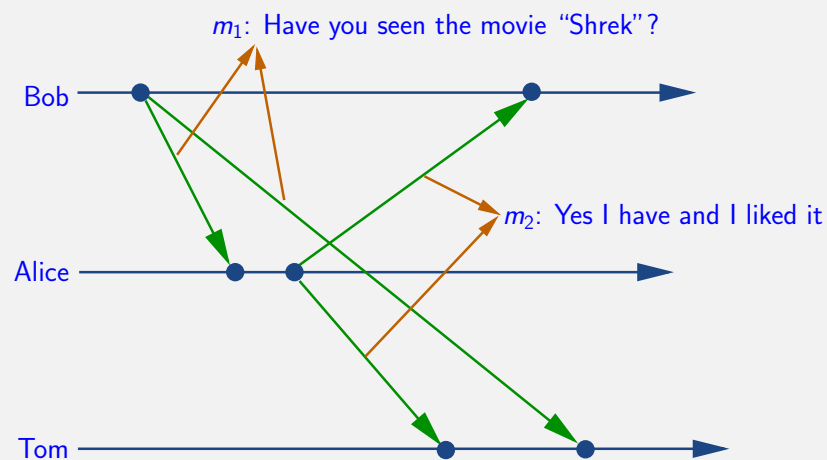
Ordering of Messages

State of a Distributed System

Monitoring a Distributed System

Ordering of Messages


- ▶ For many applications, messages should be **delivered in certain order** to be interpreted meaningfully
- ▶ Example:



- ▶ m_2 cannot be interpreted until m_1 has been received
- ▶ Tom receives m_2 before m_1 : an undesirable behavior

Useful Notations

- ▶ For a message m :
 - ▶ $src(m)$: source process of m
 - ▶ $dst(m)$: destination process of m
 - ▶ $snd(m)$: send event of m
 - ▶ $rcv(m)$: receive event of m



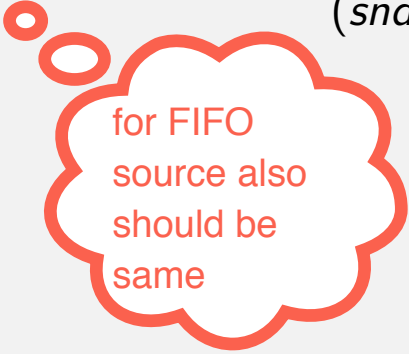
Causal Delivery of Messages

- ▶ A message w **causally precedes** a message m if $snd(w) \rightarrow snd(m)$
- ▶ An execution of a distributed system is said to be **causally ordered** if the following holds for every message m :

every message that *causally precedes* m and is *destined for the same process as* m is *delivered before* m

Mathematically, for every message w :

$$\begin{aligned} (snd(w) \rightarrow snd(m)) \wedge (dst(w) = dst(m)) \\ \implies \\ rcv(w) \rightarrow rcv(m) \end{aligned}$$



for FIFO
source also
should be
same


A Causally Ordered Delivery Protocol

- ▶ Proposed by Birman, Schiper and Stephenson (BSS)
- ▶ Assumption:
 - ▶ communication is **broadcast based**: a process sends a message to every other process
- ▶ Each process maintains a vector with one entry for each process:
 - ▶ let V_i denote the vector for process P_i
 - ▶ the j^{th} entry of V_i refers to the number of messages that have been broadcast by process P_j that P_i knows of

The BSS Protocol

- ▶ Protocol for process P_i :
 - ▶ On **broadcasting a message** m :
 - piggyback V_i on m
 - $V_i[i] := V_i[i] + 1$
 - ▶ On **arrival of a message** m from process P_j :
 - let V_m be the vector piggybacked on m
 - deliver m once $V_i \geq V_m$
 - ▶ On **delivery of a message** m sent by process P_j :
 - $V_i[j] := V_i[j] + 1$

The BSS Protocol: An Illustration




Another Causally Ordered Delivery Protocol

- ▶ Proposed by Raynal, Schiper and Toueg (RST)
- ▶ Assumption:
 - ▶ communication is **unicast based**: a process sends a message to only one other process
- ▶ Each process P_i maintains a matrix M_i of size $N \times N$
 - ▶ N is the number of processes in the system
- ▶ The entry (j, k) in the matrix M_i captures the following knowledge:
 - ▶ Suppose $M_i[j, k] = x$
 - ▶ **Interpretation**: As far as process P_i knows, process P_j has sent x messages to process P_k
- ▶ Initially, all entries in all matrices are set to 0

The RST Protocol

- ▶ Protocol for process P_i :
 - ▶ On sending message m to process P_j :
 - piggyback matrix M_i on message m
 - increment $M_i[i, j]$ by 1
 - ▶ On receiving message m from process P_j carrying matrix M_m :
 - buffer the message until $M_i[*, i] \geq M_m[*, i]$
 - ▶ On delivery of message m sent by process P_j carrying matrix M_m :
 - merge M_i and M_m
 - $\forall x, y : M_i[x, y] := \max(M_i[x, y], M_m[x, y])$
 - increment $M_i[j, i]$ by 1

The RST Protocol: An Illustration



Outline

Overview

Ordering Events

Abstract Clocks

Ordering of Messages

State of a Distributed System


Monitoring a Distributed System

State of a Distributed System

- ▶ State of a distributed system is a collection of states of all its processes and channels:
 - ▶ a process state is given by the **values of all variables** on the process
 - ▶ a channel state is given by the **set of messages in transit** in the channel
 - ▶ can be determined by examining states of the two processes it connects
- ▶ State of a process is called **local state** or **local snapshot**
- ▶ State of the system is called **global state** or **global snapshot**

Events and Local States


- ▶ A process **changes** its state by executing an event
- ▶ Example:



- ▶ Process P_1 changes its state from $x = 0$ to $x = 2$ on executing event a
- ▶ Process P_2 changes its state from $y = 2$ to $y = 3$ on executing event d

When is a Global State Meaningful?

► Example:



- Is it possible for x to be 0 and y to be 3 at the **same time**?
- Does $\{s, x\}$ form a **meaningful** global state?

Revisiting Happened-Before Relation

- ▶ Earlier, happened-before relation was defined on **events**
- ▶ The relation can be **extended to local states** as follows:
 - ▶ let s be a local state of process P_i
 - ▶ let t be a local state of process P_j
 - ▶ $s \rightarrow t$ if there exist events e and f such that:
 1. P_i executed e after s ,
 2. P_j executed f before t , and
 3. $e \rightarrow f$
- ▶ If $s \rightarrow t$, then s and t cannot **co-exist** at the same time
Also, s must occur **before** t
- ▶ If $s \parallel t$, then s and t can **co-exist** simultaneously

A Consistent Global State

- ▶ For a global state G , let $G[i]$ refer to the local state of process P_i in G
- ▶ A global state G is **meaningful or consistent** if

$$\forall i, j : i \neq j : (G[i] \rightarrow G[j]) \wedge (G[j] \rightarrow G[i])$$

$$\equiv$$

$$\forall i, j : i \neq j : G[i] \parallel G[j]$$

Recording a Consistent Global Snapshot

- ▶ Proposed by Chandy and Lamport (CL)
- ▶ Assumptions and Features:
 - ▶ channels satisfy **first-in-first-out (FIFO)** property
 - ▶ channels are **not** required to be **bidirectional**
 - ▶ communication topology may **not** be **fully connected**
- ▶ Messages exchanged by the underlying computation (whose snapshot is being recorded) are called **application messages**
- ▶ Messages exchanged by the snapshot algorithm are called **control messages**

The CL Protocol


- ▶ *Initially:* all processes are colored **blue**
- ▶ *Eventually:* all processes become **red**
- ▶ On **changing color** from blue to red:
 - record local snapshot
 - send a marker message along all outgoing channels
- ▶ On **receiving marker message** along incoming channel *C*:
 - if color is blue then
 - change color from blue to red
 - endif
 - record state of channel *C* as application messages received along *C* since turning red

The CL Protocol (Continued)


- ▶ Any process can initiate the snapshot protocol by **spontaneously changing its color** from blue to red
 - ▶ there can be multiple initiators of the snapshot protocol
- ▶ **Global Snapshot:** local snapshots of processes *just after they turn red*
- ▶ **In-Transit Messages:** *blue* application messages received by processes *after they have turned red*
- ▶ Why is the global snapshot consistent?
 - ▶ Assume an application message has the same color as its sender
 - ▶ Can a blue process receive a red application message?

The CL Protocol: An Illustration

- ▶ Three processes: P_1 , P_2 and P_3
- ▶ Five channels: $C_{1,2}$, $C_{1,3}$, $C_{2,1}$, $C_{2,3}$ and $C_{3,1}$



The CL Protocol: An Illustration (Continued)



Outline

Overview

Ordering Events

Abstract Clocks

Ordering of Messages


State of a Distributed System

Monitoring a Distributed System

A Subclass of Distributed Computation

- ▶ Many distributed computations obey the following paradigm:
 - ▶ A process is either in **active** state or **passive** state
 - ▶ A process can send an application message *only when it is active*
 - ▶ An active process can become passive at any time
 - ▶ A passive process, on receiving an application message, becomes active
- ▶ Intuitively:
 - ▶ If a process is active, then it is doing some work
 - ▶ If process is passive, then it is idle
 - ▶ An active process uses an application message to send a part of its work to another process

Distributed Computation: An Illustration



Termination Detection

- ▶ To detect if the computation has finished doing all the work
 - ▶ all processes have become passive, and
 - ▶ all channels have become empty
- ▶ Different types of computations:
 - ▶ **diffusing**: only one process is active in the beginning
 - ▶ **non-diffusing**: any subset of processes can be active in the beginning
 - ▶ no process knows which processes are active and which processes are passive

Huang's Protocol

- ▶ *Assumption:* computation is diffusing
 - ▶ the initially active process is called the **coordinator**
 - ▶ coordinator is responsible for detecting termination
- ▶ *Initially:*
 - ▶ coordinator has a weight of 1
 - ▶ all other processes have a weight of 0
- ▶ *Invariants:*
 - ▶ total amount of weight in the system is 1
 - ▶ a non-coordinator process has a non-zero weight *if and only if* it is active
 - ▶ a channel has a non-zero weight *if and only if* it is non-empty
 - ▶ weight of a channel is the sum of the weight of all messages in it

Huang's Protocol (Continued)

- ▶ Actions:
 - ▶ On **sending** an application message:
 - ▶ send half of the weight along with the message
 - ▶ On **receiving** an application message:
 - ▶ add the weight of the message to the current weight
 - ▶ On **becoming** passive:
 - ▶ send the current weight to the coordinator
- ▶ Coordinator announces termination once:
 - ▶ it has become passive and
 - ▶ it has collected all the weight

Huang's Protocol: An Illustration

